

# AffordMed URL Shortener — Professional System Design

Version: 1.0

## 1. Executive Summary

AffordMed is a lightweight URL shortening solution originally implemented as a frontend-only application (React + Vite + TailwindCSS). This document provides a professional system design for both the existing frontend-only implementation and an extended production-ready architecture that includes a backend API, database, analytics, and deployment recommendations. The goal is to provide a clear blueprint for developers and stakeholders to implement, scale, and operate the service.

## 2. Goals & Non-Functional Requirements

**Functional Goals:** shorten long URLs, redirect users, track click events, show statistics. **Availability:** 99.9% uptime for redirects (production architecture). **Latency:** redirect latency < 100ms for cached requests. **Scalability:** support >10M shortened links and 100k redirects/sec with horizontal scaling. **Security:** protect against open redirect misuse, XSS, and injection attacks. **Durability:** data persisted in durable storage with backups.

## 3. Current (Frontend-only) Architecture

The current implementation is a single-page application (SPA). All logic runs in the browser: - URL validation, short-code generation (nanoid), storage (localStorage), redirect handling via React Router. - Event logging is implemented in sessionStorage for local analytics. Limitations: links remain local to the browser; not shareable across devices; storage clearing removes data.

## 4. Proposed Production Architecture

To make AffordMed a production-ready, shareable service, introduce a backend API and managed database. Key components: - Load Balancer (e.g., AWS ALB) - API Servers (Node.js + Express) - Primary Database (MongoDB / DynamoDB) - Cache Layer (Redis) for hot redirect keys - Object store for static assets (S3) and image hosting - Analytics pipeline (Kafka -> Clickstore) for click processing - Monitoring & Logging (Prometheus, Grafana, ELK) - CI/CD pipeline (GitHub Actions -> Build -> Deploy)

## 5. Component Design

| Component        | Responsibility  | Technology                   |
|------------------|---|------------------------------|
| Web Client (SPA) | Collect URLs, show stats, initiate shorten & redirect                 | React, Vite, TailwindCSS     |
| API Server       | Validate input, generate short codes, persist mappings, expires, type | Node.js, Express, TypeScript |
| Database         | Store mappings & metadata (createdAt, expiresAt, type)                | MongoDB / DynamoDB           |
| Cache            | Serve high-frequency redirect lookups                                 | Redis / ElastiCache          |
| Message Queue    | Buffer click events for async processing                              | Kafka / AWS SQS              |
| Analytics Store  | Store aggregated click metrics  | Clickhouse / BigQuery        |

## 6. Data Model

Primary collection/table: **links** Fields (MongoDB-like): - \_id (ObjectId) - shortCode (string, unique, indexed) - longUrl (string) - createdAt (ISODate) - expiresAt (ISODate, optional) - ownerId (string, optional) — for authenticated users - clicksCount (number) - meta { title, tags, description } - disabled (boolean)

Secondary collection: **clicks** - `_id` - `shortCode` (string, indexed) - `clickedAt` (ISODate) - `referrer` (string) - `userAgent` (string) - `ipAddress` (string) - `geo` { `country`, `region`, `city` }

## 7. API Design (REST)

1. POST `/api/shorten` - Request: { `longUrl`, `customAlias` (optional), `expiryMinutes` (optional) } - Response: { `shortUrl`, `shortCode`, `createdAt`, `expiresAt` } 2. GET `/:shortCode` - Behavior: lookup `shortCode` in Redis cache -> DB -> if found increment click (async) and redirect (302) to `longUrl`. - If not found -> 404 page 3. GET `/api/link/:shortCode` - Response: link metadata (`longUrl`, `createdAt`, `expiresAt`, `clicksCount`) 4. GET `/api/stats/:shortCode` - Response: aggregated click stats (`last24h`, `last7d`), `top-referrers`

## 8. Sequence Diagrams (High-level)

Shorten flow: User -> Web Client -> POST `/api/shorten` -> API Server validates -> DB save -> return `shortCode` -> display to user  
Redirect flow: User -> GET `/:shortCode` -> Load Balancer -> API Server checks Redis -> hit -> return `longUrl` (302) -> async enqueue click event for analytics

## 9. Scaling & Caching Strategy

- Use Redis as primary cache for redirect lookups. Populate cache on first read and via write-through for new short codes.
- Partition database by sharding on `shortCode` or use consistent hashing for distributed keys.
- Auto-scale API servers based on CPU/Request metrics.
- Use CDN (CloudFront) for static assets and assets hosting.

## 10. Security & Abuse Mitigation

- Validate `longUrl` to prevent SSRF / open redirects to internal networks.
- Rate limit endpoint (POST `/api/shorten`) per IP or per API key.
- CAPTCHA for anonymous users after threshold reached.
- Store IP-blocklist and detect suspicious patterns (mass creation, repeated redirects).
- Sanitize outputs to avoid XSS in any preview fields.
- Use HTTPS everywhere and set secure cookies for auth.

## 11. Monitoring & Observability

- Metrics: request rates, error rates, latency, redirect throughput.
- Logs: structured logging for API requests and redirect events.
- Traces: distributed tracing (Jaeger) for request flows in production.
- Alerts: P95 latency > Xms, error rate > Y% trigger alerts.

## 12. Deployment Recommendations

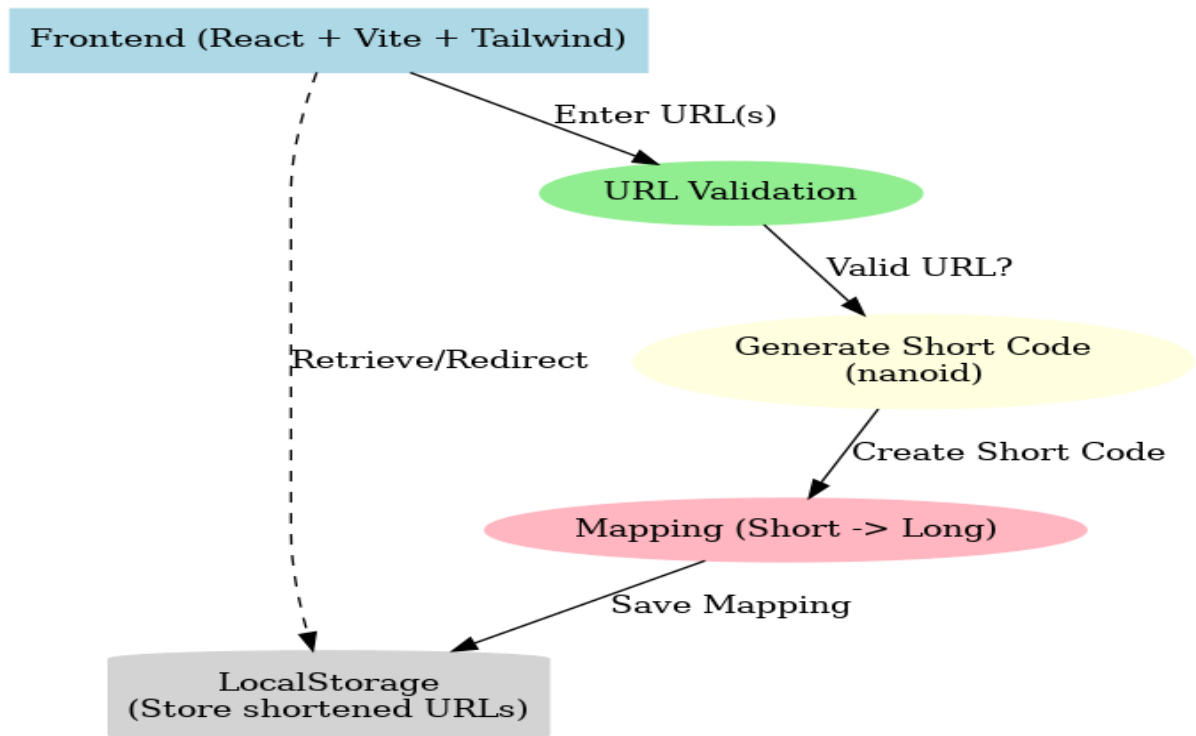
- Use containerized services (Docker). Deploy via Kubernetes or managed services (ECS/EKS).
- CI/CD: GitHub Actions -> Build -> Run tests -> Deploy to staging -> Canary -> Production.
- Use managed databases (MongoDB Atlas) and managed Redis.
- Backups: regular backups for DB; snapshots for persistent volumes.

## 13. Cost Estimate (Monthly Ballpark)

- Small MVP: \$50–\$200 (single small VPS + managed DB)
- Production (medium): \$500–\$2,000 (Load balancer, multiple app servers, managed DB, Redis, monitoring)
- Large scale: \$5k+/month (auto-scaling, analytics, multi-region)

## 14. Trade-offs & Alternatives

- Using serverless (AWS Lambda + DynamoDB) reduces ops but increases cold-start risk for redirects.
- Using SQL DB can provide strong consistency, but NoSQL (Mongo/Dynamo) offers flexible schema for metadata.
- Caching is essential for low-latency redirects; consider edge caches (Cloudflare Workers) for global performance.



## Appendix A — Data Schema (Example Documents)

Example **links** document: {"shortCode":"abc123", "longUrl":"https://example.com", "createdAt":"2025-08-26T13:00:00Z", "expiresAt":"2025-08-26T13:30:00Z", "clicksCount":0, "ownerId":null} Example **clicks** document: {"shortCode":"abc123", "clickedAt":"2025-08-26T13:05:00Z", "referrer":"https://google.com", "userAgent":"...", "ipAddress":"1.2.3.4"}

Document prepared by: AffordMed Engineering